

THE PRACTICAL PLAYBOOK

Building an AI Influencer

A step-by-step guide to designing a fictional persona, keeping their face consistent across hundreds of images, producing realistic photo and video content, and growing an audience responsibly.

10 steps

From blank page to first post

2 routes

Home PC or cloud, no GPU needed

Checklists

Templates, QC and weekly routine

What's inside

Work through the guide in order the first time. Part 2 is the part people skip, and it's the reason most AI influencers look like a different person in every post.

PART 1 Foundations

00	Ground rules before you start	4
—	The pipeline at a glance	5
01	Design the persona	6
02	Choose your tools and set up	7

PART 2 Locking the identity

03	Casting: find the face	9
04	Build the reference dataset	9
05	Train the persona LoRA	11
06	The consistency test	12

PART 3 Making content

07	Prompting for realism	14
08	From photo to video	15
09	Quality control and export	16

PART 4 Growing the account

10	Launch, post and grow	18
—	Monetisation options	19
—	The weekly routine	20

EXTRAS Reference

A	Troubleshooting	21
B	Glossary	22
C	Launch checklist	23

HOW TO READ THIS GUIDE

Each step opens with three facts: **how long it takes**, **what you need** and **what you'll have at the end**. If you can't produce the output, don't move on. Every later step depends on it.

01

PART ONE

Foundations

The rules that keep the account alive, the shape of the whole workflow, and the one document every image will come from.

Ground rules before you start

Step 00

The pipeline at a glance

Overview

Design the persona

Step 01

Choose your tools and set up

Step 02

Ground rules before you start

An AI influencer is a fictional person whose photos and videos are generated, not shot. These rules come first because breaking any of them can end the account overnight, and some can get you into legal trouble.



100% fictional

No real person's face, name or photos as a base. Never prompt "looks like <celebrity>". Never use reference photos of people who haven't signed a release.

21+

Clearly adult

State an age of 21 or older and keep adult styling and contexts. Never use youthful cues such as "teen", "girl", "petite" or school uniforms, not even as a joke.



Label it as AI

Turn on each platform's AI-content toggle and say "AI creator" in the bio. Most platforms now require this, and undisclosed AI personas get removed.



Keep a generation log

Record the prompt, model, seed, date and persona for every image you keep. If someone disputes a post, the log proves it's synthetic and original.

Two content tiers

Decide now what goes where. Mainstream platforms don't allow nudity, and posting the wrong tier is the fastest way to lose an account.


TIER	WHAT IT IS	WHERE IT CAN GO
Public	Lifestyle, fashion, fitness, travel. Flirty is fine; nudity is not.	Instagram, TikTok, Threads, YouTube, Facebook, X
Restricted	Anything more revealing than a mainstream platform allows.	Only platforms that explicitly allow it, with age verification and their AI-creator rules followed. Never on mainstream social.

HARD LINES

Deepfake and non-consensual imagery laws (e.g. the US TAKE IT DOWN Act, the UK Online Safety Act and the EU AI Act's transparency rules) target exactly the mistakes above. Treat them as non-negotiable. This guide isn't legal advice; check the rules where you live and on every platform you use.

The pipeline at a glance

There are two loops. You build the persona **once**, carefully. Then you run a cheap, fast production loop **every week**. Almost all quality problems come from rushing the first loop.



Why the identity lock matters

A text prompt alone invents a slightly different face on every render. Your audience follows a *person*, so a face that drifts between posts kills trust immediately. The fix has two parts:

The dataset (Step 04)

Shows the model one face in many situations: different angles, lighting, outfits and places.

The LoRA (Step 05)

A small add-on model that learns that face and recalls it whenever you use a made-up trigger word.

Time and money, roughly

~10 h

to build and lock the persona the first time

4-6 h

per week to produce and post 7 pieces

~\$1-5

to rent a cloud GPU for one LoRA training run

15x

a video costs about 15x an image, so iterate on stills

Estimates for a beginner on the home-PC route. Hosted tools trade time for a monthly fee; check current pricing before you commit.

Design the persona

The persona bible is the single source of truth. Every image prompt and every caption is built from it. Write it once, then **never change the physical description**.

TIME

1–2 hours

YOU NEED

A niche and a notebook

OUTPUT

A one-page persona bible

Start with the niche, not the face

People follow an account for what it gives them: outfit ideas, gym motivation, travel inspiration, a sense of humour. Pick a niche you can make 7 posts a week about for a year, then design a person who'd live in it.

Example persona bible



Ava Lane

26 · fictional · AI creator

NICHE
Everyday lifestyle, fitness and weekend travel

TRIGGER WORD
avlx woman

VOICE
Warm, playful, a bit sarcastic

IDENTITY (NEVER CHANGES)
26-year-old woman, oval face, high cheekbones, hazel-green eyes, light freckles across the nose, long wavy chestnut-brown hair with soft curtain bangs, small beauty mark above left lip, fair warm skin, slim athletic build.

PHOTO STYLE
Raw, unedited phone photo; candid; natural skin texture with visible pores; natural, uneven light; slight sensor noise.

PERSONALITY
Morning-coffee person, gym four times a week, weekend trips. Self-deprecating humour.

WARDROBE
oversized knits linen sets slip dresses athleisure denim

RECURRING LOCATIONS
sunlit apartment café window seat gym beach boardwalk
rooftop at dusk

CONTENT PILLARS
get ready with me coffee & mornings gym travel golden hour

Write a distinctive identity

Do

- ✓ Use concrete anchors: eye colour, exact hair colour, length and cut, build
- ✓ Add one or two signature marks, like freckles or a beauty mark
- ✓ Invent a rare trigger token (e.g. **avlx**) plus "woman" or "man"
- ✓ Fix 3–5 recurring locations so the world feels real

Avoid

- ✗ Vague words like "beautiful" or "stunning", which produce the generic AI face
- ✗ Real names, celebrity references or "looks like..."
- ✗ Changing hair colour or features later, which breaks the identity lock
- ✗ Any youthful descriptor, ever

Choose your tools and set up

There are two ways to do this. The home-PC route is free per image and gives you full control. The cloud route needs no graphics card and gets you started in an afternoon. The steps in this guide work for both.

TIME	YOU NEED	OUTPUT
1 hour (PC) · 15 min (cloud)	A GPU with 8 GB+, or a browser	A working image generator

Route A: home PC

Best value and most control once it's set up.

- **App:** ComfyUI (Windows portable, NVIDIA build)
- **Model:** an SDXL photoreal checkpoint, e.g. RealVisXL V5.0 or Juggernaut XL
- **Add-ons:** ComfyUI-Manager, IPAdapter Plus, Impact Pack (FaceDetailer)
- **Cost per image:** electricity only

Route B: cloud, no GPU

Fastest start. You pay per credit or per month.

- **Option 1:** rent a GPU by the hour (RunPod, Vast.ai) and run ComfyUI there
- **Option 2:** a hosted generator with a "character" or face-reference feature
- **Check:** commercial-use rights, content policy, export without watermark

Hardware for the home-PC route

PART	MINIMUM	COMFORTABLE	WHY IT MATTERS
GPU (NVIDIA)	8 GB VRAM (e.g. RTX 2070/3060)	12–24 GB	Sets which models fit and how fast images render
System RAM	16 GB + pagefile on SSD	32 GB	Model loading spikes RAM; without a pagefile, ComfyUI crashes
Storage	100 GB free SSD	250 GB+	Each checkpoint is 6–7 GB; renders add up quickly

Setting up ComfyUI

1 Download and extract

Get *ComfyUI Windows portable (NVIDIA)* from the ComfyUI GitHub releases and extract it to a short path on an SSD, e.g. `C:\AI\ComfyUI`.

2 Start it

Run `run_nvidia_gpu.bat`. The interface opens in your browser at `127.0.0.1:8188`. If you see out-of-memory errors, add `--lowvram` to the launch line.

3 Add a checkpoint

Put your SDXL model in `models\checkpoints\`. A "Lightning" variant renders in 6 steps (10–20 s per image on 8 GB); the full model is slower but has more natural skin. On 8 GB cards stick to SDXL: Flux runs, but at 1.5–3 minutes per image.

4 Install the Manager

If there's no *Manager* button, install ComfyUI-Manager. You'll use it to add IPAdapter in Step 04.

02

PART TWO

Locking the identity

Turn a description into one recognisable face, then teach a model to reproduce it on demand. This is the part that separates a believable influencer from a stream of strangers.

Casting: find the face

Step 03

Build the reference dataset

Step 04

Train the persona LoRA

Step 05

The consistency test

Step 06

Casting: find the face

Generate a batch of simple headshots from the identity description alone, then choose **one** hero face. This is the cheapest moment to change your mind, so take your time.

TIME

~30 minutes

YOU NEED

Persona bible + generator

OUTPUT

One hero headshot

1 Render about 40 headshots

Use the identity plus a neutral setup: front-facing, plain background, soft even light. Change only the seed between images.

2 Shortlist five

Look for a face that's attractive but *specific*: something you'd recognise again. Generic "model faces" are hard to lock.

3 Pick one

It should be clear, front-facing, evenly lit, with clean eyes and teeth and nothing odd. Save it as your hero image (e.g. `hero_ava.png`).

4 Not happy?

Edit the identity text, not the seed, and render again. Every face after this point is built from the hero.

CASTING PROMPT

close-up portrait photo of a 26 year old woman, oval face, high cheekbones, hazel green eyes, light freckles across the nose, long wavy chestnut brown hair with soft curtain bangs, small beauty mark above left lip, looking at camera, neutral expression, plain light grey wall, soft window light, raw unedited phone photo, natural skin texture with visible pores

Negative: smooth skin, beauty filter, airbrushed, heavy makeup, hdr, oversaturated, cartoon, illustration

PICK WITH YOUR PHONE

Look at the shortlist on your phone at Instagram size. That's how your audience will see the face, and odd features are easier to spot there than on a big monitor.

Build the reference dataset

Now produce the same face in many different situations. You'll use the hero image as a *visual* reference through IPAdapter, a node that steers each render toward a given face.

TIME

2–3 hours

YOU NEED

Hero image + IPAdapter

OUTPUT

25–40 curated images

1 Install IPAdapter

In ComfyUI-Manager, install *ComfyUI_IPAdapter_plus* and restart. Download the SDXL PLUS FACE (portraits) model into `models\ipadapter\` and the ViT-H CLIP vision model into `models\clip_vision\`.

2 Wire the reference

Load Checkpoint → IPAdapter Unified Loader → IPAdapter Advanced (with your hero image) → KSampler.
Start with weight **0.65** and end_at **0.8** so the last steps can draw natural skin.

3 Render about 45 varied images

Vary one thing at a time: shot type, angle, lighting, outfit, location, expression.

4 Curate hard

Keep 25–40. Whatever you keep gets baked into the persona permanently.

Coverage to aim for

CATEGORY	INCLUDE
Framing	Close-up (~40%), head and shoulders, waist-up, full body (~20%)
Angles	Front, ¾ left, ¾ right, profile, slightly above, slightly below
Light	Window daylight, golden hour, overcast outdoor, warm indoor lamp, flash at night
Context	3–5 outfits and 4–6 backgrounds from the bible; a few different expressions

Keep the image if

- ✓ It's unmistakably the same face
- ✓ Eyes, teeth, hands and ears are clean
- ✓ It adds something new: angle, light or setting

Reject the image if

- ✗ It's near-identical to one you've kept
- ✗ The skin looks waxy or airbrushed
- ✗ There's any distortion, even a small one

FIX REALISM NOW, NOT LATER

The LoRA learns whatever look the dataset has. A plastic-looking dataset gives you a plastic-looking persona forever. If images look "too AI", read Step 07 before you train.

Train the persona LoRA

A LoRA is a small add-on file (tens of megabytes) that teaches your checkpoint one new concept: your persona's face, recalled by the trigger word. It's what makes post #300 look like post #1.

TIME

1 h (cloud) · 3–5 h (8 GB PC)

YOU NEED

Curated dataset + captions

OUTPUT

persona_v1.safetensors

1. Prepare captions

Give each image a text file with the same name. Start with the trigger, then describe the *scene*, not the face. That teaches the model that only the person belongs to the trigger word, while clothes and backgrounds stay flexible.

EXAMPLE CAPTION · IMG_014.TXT

avlx woman, sitting at a café window seat holding a coffee cup, cream knit sweater, overcast daylight, waist-up, candid

Use folder naming like `img\10_avlx woman\` for kohya-style trainers, where 10 is the number of repeats per epoch.

2. Choose where to train

On your PC

Use kohya_ss (bmaltais GUI), OneTrainer or AI-Toolkit. Possible on 8 GB with the settings below. Close everything else, and expect a few hours.

Rent a GPU

recommended

On RunPod or Vast.ai, start an RTX 4090 with a kohya_ss template. Upload the dataset, train, and download the file. Usually about an hour and a few dollars.

3. Settings that work for SDXL on 8 GB

SETTING	VALUE	NOTE
Base model	Your SDXL checkpoint	Or base SDXL 1.0 for wider compatibility
Network dim / alpha	16 / 8	Enough capacity for one face
Resolution	1024, bucketing on	Mixed aspect ratios are fine
Batch size	1	Raise it on bigger GPUs
Optimizer · learning rate	Adafactor · 1e-4	relative_step off
Train text encoder	Off	UNet only saves memory
Mixed precision	fp16	Use bf16 on RTX 30-series and newer
Gradient checkpointing · cache latents	On · on (to disk)	Both are essential on 8 GB
Total steps	3,000–4,000	≈ images × repeats × epochs; save every 2 epochs

SAVE CHECKPOINTS ALONG THE WAY

Saving every couple of epochs gives you several versions to compare in Step 06. The last one isn't always the best. Over-trained LoRAs copy the dataset's outfits and backgrounds.

STEP 06

Identity

The consistency test

Before making a single post, prove the identity holds. Render the same 10 very different scenes with fixed seeds, and check whether they show the same person.

TIME	YOU NEED	OUTPUT
20 minutes per version	LoRA + 10 test prompts	A pass, or a clear fix

Use scenes the dataset *didn't* contain: a night market, a ski slope, a library, a kitchen, a car interior. Fixed seeds mean you can compare LoRA versions or strengths fairly, because only the LoRA changes.

RESULT	WHAT TO DO
Same face in 8+ of 10	Pass. Lock this LoRA version and strength in the persona bible.
Face drifts between scenes	Raise LoRA strength toward 1.0, or retrain with more close-ups.
Every image looks like a dataset photo	Lower strength to ~0.7, use an earlier epoch, or retrain with more variety.
Looks right but plastic	Go to Step 07. A realism pass fixes most of this without retraining.

03

PART THREE

Making content

Write prompts that look like real phone photos, turn the best stills into short videos, and catch the flaws before your audience does.

Prompting for realism

Step 07

From photo to video

Step 08

Quality control and export

Step 09

Prompting for realism

The goal isn't a perfect photo. It's a believable one. Real influencer content is shot on phones, in real places, with imperfect light. Prompt for that.

The prompt formula

trigger + shot type + action / pose + outfit + location & light + photo style

EXAMPLE · MONDAY COFFEE POST

avlx woman, medium shot, sitting on a windowsill holding a coffee mug, oversized white shirt, messy hair, small apartment with plants, morning sunlight, raw unedited iPhone photo, candid and unposed, natural skin texture with visible pores, slight sensor noise

Negative: smooth skin, beauty filter, glossy skin, porcelain skin, oversaturated, hdr, bokeh, perfect symmetry, airbrushed

Words that add realism

raw unedited phone photo candid visible pores
natural uneven lighting slight sensor noise
real-world surroundings fine facial hair

Words that scream "AI"

masterpiece 8k ultra detailed
instagram aesthetic bokeh studio lighting
flawless skin

Fixing the "obviously AI" look

SYMPTOM	USUAL CAUSE	FIX
Waxy, airbrushed skin	Face reference active until the last steps	IPAdapter end_at 0.8, weight 0.6–0.65; add a hires pass without the reference
Mushy face in full-body shots	Face too small at base resolution	Hires pass at 1.5× (denoise 0.35–0.4), then FaceDetailer
Glossy "photoshoot" look	Prompt words like <i>studio</i> , <i>bokeh</i> , <i>8k</i>	Swap them for phone-photo wording
Over-saturated, harsh contrast	CFG too high	Lightning models: CFG 1–1.5. Full models: 3.5–5
Face looks pasted on	Reference weight too high	Lower IPAdapter weight to 0.5–0.6
Everything too clean	No imperfections	Add subtle grain on export; mess up the scene a little

BIGGEST SINGLE UPGRADE

FaceDetailer (from ComfyUI-Impact-Pack) finds the face, re-renders it at high resolution and blends it back. It's the best fix for waist-up and full-body shots. Optionally add a realism LoRA ("amateur photo", "skin detail") at 0.3–0.6 strength.

From photo to video

Short vertical video is where reach lives. Video is expensive, so the rule is simple: **iterate on cheap stills, and animate only the frames you've already approved.**

TIME ~1 hour per week	YOU NEED Approved stills + a video model	OUTPUT 5–10 s clips
--------------------------	---	------------------------

Always image-to-video

Start every clip from an approved still of your persona. Text-to-video invents a new face; image-to-video keeps yours. Generate 4–8 candidate first frames per idea, pick the best one, then animate it.

MOTION PROMPT · KEEP IT TO ONE ACTION

slow push-in, she glances up from her coffee and smiles at the camera, hair moves slightly in the breeze, natural handheld phone movement

Model options

MODEL	STRENGTHS	USE IT FOR
Kling	Strong human motion and faces	A good default for persona clips
Veo	The most cinematic and realistic	Hero clips. Often the strictest content filters
Hailuo / MiniMax, Seedance	Cheaper and faster	B-roll, loops, quick tests
Wan (open weights)	Runs locally with enough VRAM	Unlimited iteration if your GPU can handle it

Run a bake-off once: put the same still through each model with the same motion prompt, then standardise on the winner. Models update often, so repeat it every few months.

Building longer videos

- Clips come out at about 5–10 seconds. Chain them by using the **last frame** of clip 1 as the **first frame** of clip 2.
- Cut 2–4 clips into a 7–20 second short, vertical 9:16 at 1080×1920.
- Put a text hook in the first second; add captions and audio *inside each app* to stay on licensed music.
- For a voice, pick one text-to-speech voice and never change it. Voice is part of the identity too.

Quality control and export

One bad hand in one post gets screenshotted and becomes the comment section. A two-minute check per image is the cheapest insurance you have.

QC checklist: reject if any fail

- ☐ Hands and fingers: count them
- ☐ Eyes (matching, same direction), teeth, earrings
- ☐ Face matches the persona; no drift
- ☐ No garbled text, signs or logos
- ☐ Straight lines stay straight (doors, tiles, railings)
- ☐ Video: no flicker, morphing or melting hands
- ☐ Nothing that could read as underage
- ☐ Right tier for the platform it's going to

Fixing small defects

Don't throw away a great image for one bad finger. Use **inpainting**: mask the problem area and regenerate only that part with the same prompt.

For a sharper final version of a winning draft, re-render it with the same seed plus a hires pass.

Export rules

- Export as JPEG and **strip all metadata**. Generator PNGs contain your full prompts and workflow.
- Add subtle grain to posts only, never to training images.

Platform sizes

WHERE	SIZE	RATIO	NOTES
Instagram / Threads feed	1080 × 1350	4:5	Takes up the most screen in the feed
Stories, Reels, TikTok, Shorts	1080 × 1920	9:16	Keep text out of the top and bottom 15%
X	1200 × 1500	4:5	Flag media as sensitive where required

04

PART FOUR

Growing the account

Set up accounts that survive, warm them up, post consistently, read the numbers, and turn attention into income.

Launch, post and grow

Step 10

Monetisation options

Income

The weekly routine

Routine

Launch, post and grow

New accounts are watched closely by platforms. A slow, human-looking start protects your reach for months.

Account setup

- ☐ New email address and phone number per persona
- ☐ Instagram **Creator** account (unlocks insights and scheduling)
- ☐ Bio states the persona is an AI creator
- ☐ Profile photo: your best close-up from the dataset
- ☐ 9–12 posts ready before you start promoting

 **ava.lane**

AI info

ava.lane Third coffee and it's not even 10am ☕ Be honest, how many are you on? [#morningroutine](#) [#slowliving](#) [#coffeetime](#)

Warm-up plan

WHEN	POSTING	ACTIVITY
Week 1	One post every 1–2 days	15–20 min a day engaging by hand in your niche. No links.
Week 2	Daily posts + 3–5 stories	Keep engaging. Still no links.
Week 3+	Daily posts, 2–4 Reels a week	Add a link in the bio. Start collaborations.

Captions in the persona's voice

- **Line 1:** a hook, the first thing people read
- **Line 2:** a tiny story or detail that makes them feel real
- **Line 3:** a question that invites replies
- 3–5 niche hashtags, not 30. Reply to early comments in character.

What to measure

Saves

The strongest signal that a post is worth something

Shares

What pushes you to non-followers

Profile visits

Curiosity about the persona

Follows / post

The number that actually grows the account

Each week, copy the scenes, outfits and hooks that won and drop the ones that didn't. Schedule posts with Meta Business Suite (Instagram and Facebook, free) and X's built-in scheduler.

Monetisation options

Monetise once you have consistent engagement, not just followers. A few thousand engaged followers in a clear niche are worth more than 50,000 random ones.

1

Brand partnerships

Paid posts featuring products in your niche. Disclose them as ads ([#ad](#) or the paid-partnership tag), and tell the brand up front that the creator is AI.

2

Affiliate links

A commission on products the persona "uses": outfits, gym gear, coffee kit. Works from week 3 with a link-in-bio page.

3

Digital products

Photo presets, wallpapers, guides, prompt packs or a newsletter. You set the price and don't depend on a sponsor.

4

Subscriptions

Exclusive content on a platform that explicitly allows AI creators. Follow its labelling and age-verification rules.

5

Platform payouts

Creator funds and ad-revenue programmes. Small at first, but they grow with reach. Check each programme's AI policy.

6

Services

Once it works, sell the skill: build personas or content for brands that want their own AI ambassador.

READ THE SMALL PRINT

Payment processors, ad networks and platforms each have their own rules on AI-generated people. Read them before you build income on one, and spread your income across at least two sources.

The weekly routine

Consistency beats intensity. This routine fits around a job: about 4–6 hours of production a week plus 15–20 minutes of engagement a day.

MON Review last week's numbers. Write 7 scenes from your pillars.	TUE Render 4 drafts per scene with the fast preset.	WED Pick winners, re-render sharp, QC, inpaint fixes.	THU Animate 2–4 stills. Edit Reels.	FRI Export, write captions, schedule the week.	SAT Stories and replies in character.	SUN Light engagement. Log what worked.
--	---	---	---	--	---	--

Example shot list for one week

DAY	PILLAR	SCENE PROMPT (ADDED AFTER THE TRIGGER)
Mon	Coffee & mornings	sitting on a windowsill with a coffee mug, oversized white shirt, messy hair, morning sunlight
Tue	Gym	mirror selfie at the gym, sage green sports set, ponytail, slight smirk
Wed	Get ready with me	sitting at a vanity putting on earrings, black satin dress, warm bedroom with fairy lights
Thu	City	full body crossing the street, beige trench coat, sunglasses, autumn city
Fri	Golden hour	rooftop at golden hour, linen sundress, wind in hair, looking back over shoulder
Sat	Travel	walking on a beach boardwalk, open linen shirt, straw hat, bright midday sun
Sun	Cozy	lying on a bed reading a book, oversized knit sweater, rainy window, soft lamp light

Troubleshooting

The problems almost everyone hits in the first two weeks, and what to do about them.

PROBLEM	FIX
CUDA out of memory	Add <code>--lowvram</code> to the launch file. Use a Lightning checkpoint. Close other GPU apps, including games and browsers with video.
ComfyUI closes while loading a model	The system ran out of RAM. Enable a pagefile on an SSD, close apps, or upgrade to 32 GB.
Images look fried or over-sharpened	Sampler settings don't match the checkpoint. Lightning: ~6 steps, CFG 1–1.5, dpmpp_sde + karras.
"Checkpoint not found"	The filename in your workflow must match the file in <code>models\checkpoints</code> exactly.
Every persona looks like the same stock model	Make the identity more distinctive: freckles, a beauty mark, a specific haircut, an unusual eye colour.
Face changes between posts	You're not using the LoRA, or it's too weak. Re-run the consistency test (Step 06).
Outfits or backgrounds repeat in every image	The LoRA is over-trained. Use an earlier epoch or lower strength, and caption scenes more fully.
Video face melts mid-clip	Use a close or medium shot as the first frame, one simple action, and a shorter clip.
Reach dropped suddenly	Check for a missed AI label, links posted too early, or a borderline post. Remove it and slow down for a week.

Glossary

Checkpoint: the main image model file (e.g. an SDXL model). Everything is generated from it.

LoRA: a small add-on file that teaches a checkpoint one new thing, like your persona's face.

Trigger word: a made-up token (e.g. `av1x woman`) that switches the LoRA's concept on in a prompt.

IPAdapter: a ComfyUI node that uses an image, such as your hero face, as a visual reference.

Seed: the random number behind an image. Same seed plus same settings gives the same image.

CFG: how strictly the model follows the prompt. Too high looks harsh and over-saturated.

Steps · sampler: how many refinement passes the model makes, and the method it uses.

Hires pass: upscale an image and redraw it lightly, which adds detail to skin and eyes.

Denoise: how much a hires or inpainting pass is allowed to change (0 = nothing, 1 = everything).

Inpainting: regenerating only a masked part of an image, like a hand or an eye.

FaceDetailer: automatically re-renders the face at high resolution and blends it back.

Image-to-video: animating an existing still, which keeps the persona's face much better than text-to-video.

VRAM: your graphics card's memory. It limits which models you can run.

Epoch: one full pass over the training dataset when training a LoRA.

Launch checklist

Tick these off in order. If one is missing, the steps after it get harder.

Foundations

- ☐ Niche chosen; it can carry 7 posts a week for a year
- ☐ Persona bible written, with an identity that won't change
- ☐ Stated age 21+; fully fictional; trigger word chosen
- ☐ Generator working (home PC or cloud)
- ☐ Generation log started

Identity

- ☐ 40 casting headshots rendered; hero face chosen
- ☐ ~45 reference images rendered with IPAdapter
- ☐ 25–40 curated, varied, realistic images kept
- ☐ Captions written: trigger + scene
- ☐ LoRA trained; consistency test passed (8+/10)

Content

- ☐ Realism prompt style saved in the bible
- ☐ First week's shot list written
- ☐ Every pick passed the QC checklist
- ☐ Exported as JPEG with metadata stripped
- ☐ Video model chosen after a bake-off

Launch

- ☐ New email and phone; Creator account
- ☐ Bio says "AI creator"; AI labels switched on
- ☐ 9–12 posts ready before promotion
- ☐ Week 1 warm-up: no links, engage daily
- ☐ Weekly review booked in the calendar

YOU'RE READY

If every box is ticked, you have something most AI influencer accounts don't: a face that stays the same, content that looks real, and an account set up to last. Now the work is consistency. Post, measure, repeat.



One face. Every post.

Build the persona once, lock the identity, then let a simple weekly routine do the rest.

VidEro